

Using BERT for Word Sense Disambiguation

Jiaju Du, Fanchao Qi, Maosong Sun

Department of Computer Science and Technology, Tsinghua University

Institute for Artificial Intelligence, Tsinghua University

State Key Lab on Intelligent Technology and Systems, Tsinghua University

{djj18, qfc17}@mails.tsinghua.edu.cn, sms@tsinghua.edu.cn

Abstract

Word Sense Disambiguation (WSD), which aims to identify the correct sense of a given polyseme, is a long-standing problem in NLP. In this paper, we propose to use BERT to extract better polyseme representations for WSD and explore several ways of combining BERT and the classifier. We also utilize sense definitions to train a unified classifier for all words, which enables the model to disambiguate unseen polysemes. Experiments show that our model achieves the state-of-the-art results on the standard English All-word WSD evaluation.

1 Introduction

Ambiguity is common in natural language. Word Sense Disambiguation (WSD) deals with lexical ambiguity, i.e. polysemes in sentences. An effective WSD tool can benefit various downstream tasks, such as Information Retrieval (Zhong and Ng, 2012) and Machine Translation (Neale et al., 2016).

Recently, the pre-trained language models, such as ELMo (Peters et al., 2018), GPT (Radford et al., 2018), and BERT (Devlin et al., 2019), have been proven to be effective to extract features from plain text. They pre-train language models on large corpora, then add the pre-trained word representation into task-specific models, or directly fine-tune the language model on downstream tasks. Peters et al. (2018) tried to incorporate the pre-trained ELMo embeddings as WSD features, but there are currently no studies which fine-tune language models on WSD task.

There have been lots of works on WSD, and they can mainly be divided into two groups: supervised methods and knowledge-based methods. Supervised methods need large sense annotated corpora. They train a classifier using features extracted from the context of the polyseme. These

methods can also be divided into two subgroups according to the features they use. Feature-based supervised methods use many conventional features like surrounding words, PoS tags of surrounding words, local word collections (Zhong and Ng, 2010), word embeddings, and PoS tag embeddings (Iacobacci et al., 2016). Neural-based supervised methods use a neural network encoder like BiLSTM to extract features (Melamud et al., 2016; Yuan et al., 2016; Raganato et al., 2017a).

Knowledge-based methods rely on the structure and content of knowledge bases, for instance, sense definitions (Lesk, 1986; Basile et al., 2014) and semantic networks which provide the relationship and similarity between two senses (Agirre and Soroa, 2009; Moro et al., 2014). Supervised methods perform better than knowledge-based methods (Raganato et al., 2017b), but knowledge-based methods are usually unsupervised and require no sense annotated data. Some recent studies have explored ways of using knowledge like sense definitions to enhance supervised methods (Luo et al., 2018b,a).

In this paper, we fine-tune BERT on the WSD task for the first time and compare the performance of different polyseme features output by the BERT encoder. The fine-tuned model beats baselines by a large margin. Many polysemes are rare, and the sense distribution of a polyseme is usually unbalanced. So the sense annotated corpora usually lack annotations for some polysemes or particular senses. To address this issue, we consider using sense definitions because we can obtain definitions for unseen sense from lexical databases easily. We find that the incorporation of sense definition improves the performance significantly.

The contributions of this paper are: (1) We fine-tune the BERT on WSD task and achieve the state-of-the-art results. (2) We prove that external knowledge is still useful for WSD with pre-trained

← optional?
lets try w/o
this first

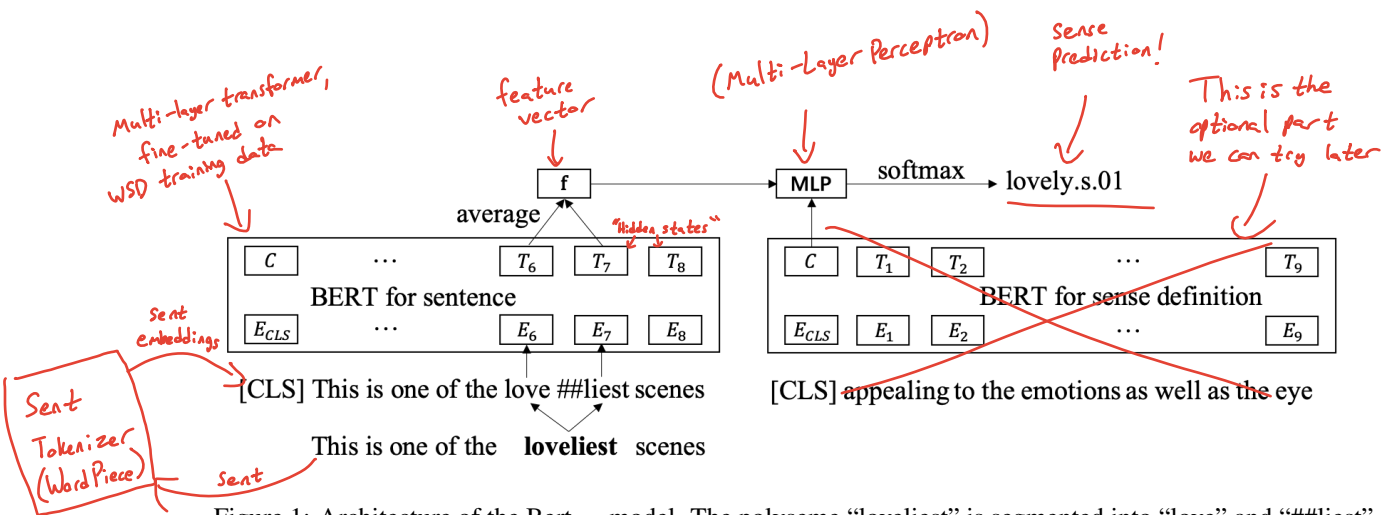


Figure 1: Architecture of the Bert_{def} model. The polyseme “loveliest” is segmented into “love” and “##liest”.

language models and improves the performance.

2 Methodology This section explains how the model works

Given a sentence and some polysems in the sentence, the WSD task aims to identify the correct senses of the polysems in the sentence. Our model consists of an encoder and a classifier. The encoder extracts the polysems’ features from sentences, and the classifier uses the features to predict senses. Figure 1 gives an overview of our model. We represent the sentence as a word sequence (w_1, \dots, w_n) , and assume that w_p is the polyseme. We denote the sense set of w_p as $S_p = \{s_1, \dots, s_{|S_p|}\}$. The definition of s_i in the lexical databases is a word sequence $d_i = (w_{i1}, \dots, w_{i|d_i|})$. The hidden size of BERT is H .

2.1 Encoder

We use BERT (Devlin et al., 2019) as the encoder. BERT uses the WordPiece (Wu et al., 2016) embeddings as a part of inputs. The WordPiece embeddings have a fixed size of vocabulary which includes some words and some word pieces, and segments out-of-vocabulary words. For example, the word “loveliest” is segmented into two word pieces “love” and “##liest”. Then a multi-layer Transformer reads the embeddings of the word piece sequence and outputs a hidden state for every token in the sequence.

As WSD is a single sentence tagging task, the most simple way of using BERT is to insert [CLS] at the start of the input sentence, and use the hidden state $h_i \in \mathbb{R}^H$ output by BERT to predict the label of the i -th token. However, we cannot directly apply this simple method on WSD because some polysems (about 15% in the training dataset) may be segmented into word pieces. These polysems correspond to at least two final hidden states. We assume that the polyseme w_p

corresponds to a hidden state list $\mathbf{h}_j, \dots, \mathbf{h}_{j+k-1}$. To obtain a fixed-sized feature vector for the classifier, we average these hidden states:

$$\mathbf{f} = \frac{1}{k} \sum_{l=0}^{k-1} \mathbf{h}_{j+l},$$

where \mathbf{f} is the feature used by the classifier. Max pooling is another way of merging these hidden states:

$$\mathbf{f} = \max_{0 \leq l \leq k-1} \mathbf{h}_{j+l}.$$

We can also get features by concatenating \mathbf{f} and the first final hidden state \mathbf{h}_0 . \mathbf{h}_0 corresponds to [CLS] and encodes global information of the sentence.

2.2 Classifier

We use a 2-layer MLP to predict the correct sense. The MLP outputs the sense distribution of the polyseme in the given context:

$$\mathbf{p} = \text{softmax}(L_2(\text{ReLU}(L_1(\mathbf{f}))),$$

where $L_i(\mathbf{x}) = \mathbf{W}_i \mathbf{x} + \mathbf{b}_i$ are fully-connected linear layers, $\mathbf{W}_1 \in \mathbb{R}^{H \times H}$, and $\mathbf{W}_2 \in \mathbb{R}^{|S_{w_p}| \times H}$. The L_2 layer is specific for every polyseme.

This MLP classifier has poor performance on unseen or infrequent words and senses because of lack of data. So we introduce the sense definitions to address the issue of data scarcity. We use another BERT to encode the definition of a sense into its sense vector. Then this sense vector is used as the parameter when predicting this sense. Formally, we replace L_2 with $L'_2(\mathbf{x}) = \mathbf{W}'_2 \mathbf{x} / \sqrt{H}$, where $\mathbf{W}'_2 = [\mathbf{d}_1; \dots; \mathbf{d}_{|S_w|}] \in \mathbb{R}^{|S_w| \times H}$ is the concatenation of all sense vectors of the polyseme. The sense vectors map senses into a unified space and encode the semantic similarities into the distance between vectors. So the classification for

Optional

	Training	Validation	Test
#Sentences	37,176	135	1,038
#Tokens	802,443	3,201	22,302
#Annotations	226,036	455	6,798
Ambiguity	6.8	8.5	5.7

Table 1: Statistics of the WSD training, validation, and test dataset. The “ambiguity” presents the average number of senses for instances in the dataset.

unseen or infrequent senses can benefit from similar senses which have more training instances. We name the model without and with definitions Bert and Bert_{def} respectively.

In the training process, parameters are updated by minimizing the cross-entropy loss between the true label y and the sense distribution p :

$$L = -\frac{1}{M} \sum_{m=1}^M \sum_{s=1}^{|S_m|} [y_m]_s \log[p_m]_s,$$

where M is the number of instances in the dataset, and y_m is a one-hot vector which represents the true label of w_m . [000100]

3 Experiments 3.1 + 3.2 explain how

3.1 Datasets to train + validate the model

We evaluate our model on the English all-words tasks. We use the evaluation framework proposed by Raganato et al. (2017b). It provides five all-words fine-grained WSD datasets for evaluation: Senseval-2 (Edmonds and Cotton, 2001, SE2), Senseval-3 task 1 (Snyder and Palmer, 2004, SE3), SemEval-07 task 17 (Pradhan et al., 2007, SE7), SemEval-13 task 12 (Navigli et al., 2013, SE13), SemEval-15 task 13 (Moro and Navigli, 2015, SE15). Following previous works, we use SE7 as the validation dataset, and use the SE2, SE3, SE13, and SE15 as test datasets. The framework provides two annotated corpora for training: SemCor (Miller et al., 1994) and OMSTI (Taghipour and Ng, 2015). We choose SemCor as our training dataset. Table 1 illustrates the statistics of these datasets. All of the datasets are annotated with WordNet (Miller, 1995) 3.0.

3.2 Experiment Setup

We use the BERT_{BASE} as the encoder. The number of Transformer layers is 12, the hidden layer size is 768, and the number of attention heads

is 12. We tried to use the BERT_{LARGE}, but the F1-score is nearly the same as BERT_{BASE}. We use Dropout in every layer of the classifier and the dropout rate is 0.5. The optimizer is Adam (Kingma and Ba, 2014). We reduce the learning rate during the training process. In the i -th epoch, the learning rate is $0.001/i$. The parameters of the BERT encoder are fixed in the first 10 epochs. We train the model for 50 epochs and choose the model which has the best F1-score on the validation set.

We compare our model with the following baselines: the simple MFS baseline which always outputs the most common sense in the training dataset, the knowledge-based methods Lesk_{+ext,emb} (Basile et al., 2014) and Babelify (Moro et al., 2014), the feature-based supervised methods IMS (Zhong and Ng, 2010) and IMS_{+emb} (Iacobacci et al., 2016), and the neural based methods Bi-LSTM (Kågebäck and Salomonsson, 2016; Raganato et al., 2017a), GAS (Luo et al., 2018b), CAN and HCAN (Luo et al., 2018a).

3.3 English All-word Task Results

Table 2 shows the F1-score of our models and baselines on the standard English All-words WSD benchmark. The Bert and Bert_{def} in the table use the average operation to merge hidden states. They don’t use the sentence vectors. Our best model improves the state-of-the-art results by 5.2%, which indicates that BERT encoder is quite powerful in the WSD task. Moreover, our models outperform previous models on all of the four datasets and PoS types. Introducing sense definitions can significantly improve the performance of the Bert model. The F1-score of Bert_{def} is better than Bert on almost all datasets and PoS types except on the adverbs, which reveals the efficiency of introducing sense definitions.

3.4 Discussion

Word Frequency We compare the performance of Bert and Bert_{def} on words with different occurrence numbers in the training dataset. Table 3 shows the results. Compared with the Bert model, Bert_{def} achieves the largest performance improvement on unseen words (8% F1-score). The improvement decreases as the word becomes frequent because the Bert model can be trained better with more instances. So we can conclude that utilizing sense definitions can enhance the model on infrequent polysemes. In addition, both models

System	Test Datasets				Concatenation of Test Datasets				All
	SE2	SE3	SE13	SE15	Noun	Verb	Adj	Adv	
MFS Baseline	65.6	66.0	63.8	67.1	67.7	49.8	73.1	80.5	65.5
Lesk _{+ext,emb}	63.0	63.7	66.2	64.6	70.0	51.1	51.7	80.6	64.2
Babelify	67.0	63.5	66.4	70.3	68.9	50.7	73.2	79.8	66.4
IMS	70.9	69.3	65.3	69.5	70.5	55.8	75.6	82.9	68.9
IMS _{+emb}	72.2	70.4	65.9	71.5	71.9	56.6	75.9	84.7	70.1
Bi-LSTM	71.1	68.4	64.8	68.3	69.5	55.9	76.2	82.4	68.4
Bi-LSTM _{+att,LEX,POS}	72.0	69.1	66.9	71.5	71.5	57.5	75.0	83.8	69.9
GAS(Concatenation)	72.1	70.2	67.0	71.8	72.1	57.2	76.0	84.4	70.3
GAS _{ext} (Concatenation)	72.2	70.5	67.2	72.6	72.2	57.7	76.6	85.0	70.6
CAN ^w	72.3	69.8	65.5	71.1	71.1	57.3	76.5	84.7	69.8
CAN ^s	72.2	70.2	69.1	72.2	73.5	56.5	76.6	80.3	70.9
HCAN	72.8	70.3	68.5	72.8	72.7	58.2	77.4	84.1	71.1
What we're replicating → Bert	74.0	73.1	71.3	74.3	75.0	61.2	77.2	86.1	73.1
The optional version → Bert _{def}	76.4	74.9	76.3	78.3	78.3	65.2	80.5	83.8	76.3

Table 2: F1-score(%) for the English all-word WSD evaluation. We report the F1-score of these systems on the four datasets (SE2, SE3, SE13, and SE15), on every part-of-speech type of polyseme (Noun, Verb, Adj, and Adv), and on the overall test dataset.

Word Count	0	1-10	11-50	51-200	>200
Bert	82.15	73.19	74.04	71.36	68.20
Bert _{def}	90.15	79.94	75.56	73.57	67.65
#Words	650	1406	1976	1854	912
Ambiguity	1.59	2.75	4.65	7.15	12.23

Table 3: F1-score(%) on words with different occurrence numbers in the training dataset.

have better performance on low-frequency words. The reason is that high-frequency words usually have many senses.

Ablation Study We compare two variations of the BERT encoder: the way of merging hidden states of polysemes, and whether concatenating the hidden states of polysemes with the sentence vector. The results of these variations are presented in Table 4. All of the models are based on Bert_{def}. We can find that the performance of using average or max operation is nearly the same for Bert_{def}. But Concatenating the sentence vector will impair the F1-score. The reason is: the sentence vector contains the global semantic information of the entire sentence. The sense of polyseme is determined by its local context. So using the sentence vector brings too much irrelevant information.

	SE2	SE3	SE13	SE15	All
Mean	76.4	74.9	76.3	78.3	76.3
Max	76.3	75.1	76.6	78.4	76.4
Mean _{Concat}	74.2	73.8	76.2	78.1	75.2
Max _{Concat}	73.4	73.7	76.1	76.7	74.8

Table 4: Ablation Study. Mean and Max denotes the average and max operation for merging hidden states of polysemes. Concat means concatenating the hidden states of polysemes with the sentence vector.

4 Conclusion

In this paper, we fine-tune the pre-trained language models like BERT on WSD tasks for the first time. We find that our BERT-based models achieve the state-of-the-art results on the standard evaluation. We also utilize sense definitions to enhance the model on infrequent polysemes. In the future works, we will consider using the relations between senses, like hypernym and hyponym, to provide more accurate sense representations.

References

- Eneko Agirre and Aitor Soroa. 2009. Personalizing pagerank for word sense disambiguation. In *Proceedings of EACL*.
- Pierpaolo Basile, Annalina Caputo, and Giovanni Se-

- meraro. 2014. An enhanced lesk word sense disambiguation algorithm through a distributional semantic model. In *Proceedings of COLING*.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL*.
- Philip Edmonds and Scott Cotton. 2001. Senseval-2: overview. In *Proceedings of SENSEVAL*.
- Ignacio Iacobacci, Mohammad Taher Pilehvar, and Roberto Navigli. 2016. Embeddings for word sense disambiguation: An evaluation study. In *Proceedings of ACL*.
- Mikael Kågebäck and Hans Salomonsson. 2016. Word sense disambiguation using a bidirectional lstm. In *Proceedings of the Workshop on Cognitive Aspects of the Lexicon*.
- Diederik P Kingma and Jimmy Ba. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- Michael Lesk. 1986. Automatic sense disambiguation using machine readable dictionaries: how to tell a pine cone from an ice cream cone. In *Proceedings of the International Conference on Systems Documentation*.
- Fuli Luo, Tianyu Liu, Zexue He, Qiaolin Xia, Zhifang Sui, and Baobao Chang. 2018a. Leveraging gloss knowledge in neural word sense disambiguation by hierarchical co-attention. In *Proceedings of EMNLP*.
- Fuli Luo, Tianyu Liu, Qiaolin Xia, Baobao Chang, and Zhifang Sui. 2018b. Incorporating glosses into neural word sense disambiguation. In *Proceedings of ACL*.
- Oren Melamud, Jacob Goldberger, and Ido Dagan. 2016. context2vec: Learning generic context embedding with bidirectional lstm. In *Proceedings of SIGNLL*.
- George A Miller. 1995. Wordnet: a lexical database for english. *Communications of the ACM*, pages 39–41.
- George A Miller, Martin Chodorow, Shari Landes, Claudia Leacock, and Robert G Thomas. 1994. Using a semantic concordance for sense identification. In *Proceedings of the workshop on Human Language Technology*.
- Andrea Moro and Roberto Navigli. 2015. Semeval-2015 task 13: Multilingual all-words sense disambiguation and entity linking. In *Proceedings of SemEval*.
- Andrea Moro, Alessandro Raganato, and Roberto Navigli. 2014. Entity linking meets word sense disambiguation: a unified approach. *Transactions of the Association for Computational Linguistics*, 2:231–244.
- Roberto Navigli, David Jurgens, and Daniele Vannella. 2013. Semeval-2013 task 12: Multilingual word sense disambiguation. In *Proceedings of SemEval*.
- Steven Neale, Luís Gomes, Eneko Agirre, Oier Lopez de Lacalle, and António Branco. 2016. Word sense-aware machine translation: Including senses as contextual features for improved translation models. In *Proceedings of LREC*.
- Matthew Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. Deep contextualized word representations. In *Proceedings of NAACL*.
- Sameer Pradhan, Edward Loper, Dmitriy Dligach, and Martha Palmer. 2007. Semeval-2007 task-17: English lexical sample, srl and all words. In *Proceedings of SemEval*.
- Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre-training.
- Alessandro Raganato, Claudio Delli Bovi, and Roberto Navigli. 2017a. Neural sequence learning models for word sense disambiguation. In *Proceedings of EMNLP*.
- Alessandro Raganato, Jose Camacho-Collados, and Roberto Navigli. 2017b. Word sense disambiguation: A unified evaluation framework and empirical comparison. In *Proceedings of EACL*.
- Benjamin Snyder and Martha Palmer. 2004. The english all-words task. In *Proceedings of SENSEVAL*.
- Kaveh Taghipour and Hwee Tou Ng. 2015. One million sense-tagged instances for word sense disambiguation and induction. In *Proceedings of CoNLL*.
- Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V Le, Mohammad Norouzi, Wolfgang Macherey, Maxim Krikun, Yuan Cao, Qin Gao, Klaus Macherey, et al. 2016. Google’s neural machine translation system: Bridging the gap between human and machine translation. *arXiv preprint arXiv:1609.08144*.
- Dayu Yuan, Julian Richardson, Ryan Doherty, Colin Evans, and Eric Altendorf. 2016. Semi-supervised word sense disambiguation with neural models. In *Proceedings of COLING*.
- Zhi Zhong and Hwee Tou Ng. 2010. It makes sense: A wide-coverage word sense disambiguation system for free text. In *Proceedings of ACL*.
- Zhi Zhong and Hwee Tou Ng. 2012. Word sense disambiguation improves information retrieval. In *Proceedings of ACL*.